

Prediction of Micromixing Scalar Concentration Field Using a Neural Network

Usama KHALID¹, Aly MORSY², Taha JAVED³, Busayomi THOMPSON-AJAYI⁴, Elin VESPER⁵, and Daniele MARCHISIO⁶

^{1,2,3,4}Politecnico di Torino, Corso Duca degli Abruzzi, 24 | 10129 Torino, Italy

⁵Bayer Aktiengesellschaft, Process Technologies, Leverkusen, Germany

⁶DISAT, Politecnico di Torino, Corso Duca degli Abruzzi, 24 | 10129 Torino, Italy

Abstract

A physics-informed convolutional neural network (CNN) is developed as a surrogate model to simulate turbulent micromixing based on high-fidelity CFD data. The network is trained on scalar transport fields generated by OpenFOAM direct numerical simulations, capturing the diffusion of a passive scalar in homogeneous isotropic turbulence. Emphasis is placed on the Batchelor-scale physics of micromixing which is the smallest scalar mixing scale and is crucial in high-Schmidt-number liquid flows. The CNN architecture (a 3D U-Net with $\sim 1.4 \times 10^6$ parameters) is designed to embed physical knowledge: it enforces non-negativity of concentration and smoothness of scalar gradients to reflect diffusive processes. This approach addresses key challenges in resolving turbulent scalar transport, including extreme resolution requirements and memory bottlenecks, by learning an efficient model that reproduces fine-scale mixing without needing an ultra-fine mesh. The study is motivated by sustainable pharmaceutical manufacturing, where improved micromixing can enhance reaction yields and reduce unwanted by-products. The trained model demonstrates good agreement with CFD benchmarks, capturing the overall mixing patterns and scalar dissipation trends across a range of Schmidt numbers. By coupling first-principles CFD data with modern deep learning, the framework achieves an order-of-magnitude speedup in predicting scalar concentration fields while maintaining physical fidelity. This work showcases how integrating convolutional neural networks with CFD can accelerate process simulations and inform the design of greener, more efficient mixing strategies in industrial reactors.

Contents

29

1	Introduction	3	30
2	Methodology	6	31
2.1	Computational Fluid Dynamics Setup	6	32
2.2	CFD Simulation Results	11	33
2.3	Neural Network Development	12	34
3	Results and Discussion	22	35
3.1	Fully Connected Network Results	22	36
3.2	Convolutional Neural Network Results	22	37
3.3	Computational Performance	25	38
4	Conclusion	26	39
5	Future Work	27	40
5.1	CFD Extensions	27	41
5.2	Neural Network Improvements	27	42
A	Summary of Solver Modifications (scalarTransportFoam)	29	43
A.1	Purpose of the Modification	29	44
A.2	Key Code-Level Changes	29	45
B	MATLAB Script Modifications for HDF5 to OpenFOAM Conversion	30	46
B.1	Purpose of the Modification	30	47
B.2	Key Code-Level Changes	31	48
C	3D CNN Implementation Code	31	49

1. Introduction

Micromixing, the molecular-scale mixing of reactive species in turbulent flows, plays a pivotal role in chemical and pharmaceutical process engineering. In industrial reactors (e.g. for pharmaceutical intermediate synthesis), incomplete micromixing allows competing side reactions to form undesired by-products, which lowers the yield of the target product [1]. Using catalysts can accelerate the primary reaction, but if mixing is slow, the reaction outcome is dictated by the competition between reaction kinetics and turbulent mixing rates [1]. In such cases, improving the micromixing efficiency directly translates to higher selectivity and less waste [1]. Consequently, understanding and predicting micromixing in turbulence is crucial for designing sustainable processes that maximize desired products while minimizing impurities.

Turbulent Scalar Transport and Micromixing Scales: Turbulent mixing occurs across a hierarchy of length scales. Energy is injected at large scales and cascades to the Kolmogorov scale (η), where viscous forces dissipate turbulence. For scalar mixing, an even smaller scale can govern: the Batchelor scale (λ_B). The Batchelor scale is the diffusion length scale below which molecular diffusion smooths out concentration fluctuations [2]. It is related to the Kolmogorov scale by $\lambda_B \approx \eta/\sqrt{Sc}$, meaning that in liquids with high Schmidt number ($Sc \gg 1$), scalar concentration gradients exist on much finer scales than the smallest eddies [3]. Figure 1 illustrates this scale separation: for example, in water, λ_B is on the order of only a few microns when η is on the order of tens of microns. These tiny mixing scales, the micromixing scales, control how quickly reactants truly molecularly mix, and hence often determine the selectivity of fast competitive reactions.

However, resolving the Batchelor scale in simulations is extremely challenging. A direct numerical simulation (DNS) must use a grid fine enough to capture λ_B , which can be several times finer than the Kolmogorov-scale grid. For instance, a DNS of homogeneous turbulence at Kolmogorov-scale resolution 1024^3 already contains over 4 billion grid points and requires on the order of 8 GB of memory per scalar field in double precision. Achieving even $2\times$ finer resolution (2048^3) to approach the Batchelor scale would raise the cell count by a factor of 8 (to $\sim 8.6 \times 10^9$ cells), and a $4\times$ refinement (4096^3) would require $\sim 6.9 \times 10^{10}$ cells – a prohibitive number of grid points and memory (hundreds of GB per snapshot). In addition, high-Schmidt-number flows demand longer simulation times to observe complete mixing, since molecular diffusion is slow. These factors result in tremendous computational expense for fully resolved micromixing studies.

Major computational challenges can be summarized as follows:

Resolution limits: Capturing Batchelor-scale gradients requires exceedingly fine grids (e.g. resolving λ_B for $Sc \gg 1$ may need billions of cells). This far exceeds the grid size of standard turbulence DNS which resolves only up to Kolmogorov scales.

Memory bottlenecks: Storing and processing 3D fields at Kolmogorov or finer resolution strains memory and I/O capacity. A single 1024^3 scalar field is ~ 4 GB (float32), and going

to finer scales multiplies this requirement. Handling many time snapshots or performing data-driven training on such fields can exhaust typical computing resources. 89
90

Schmidt number dependence: The scalar mixing problem is parameterized by Sc, which 91
affects both the required spatial resolution ($\lambda_B \propto Sc^{-1/2}$) and the mixing time scale. High-Sc 92
flows exhibit thin concentration filaments that persist longer, necessitating small time steps 93
and long integration times to capture diffusion. Each distinct Sc essentially changes the 94
physics and may even require a different grid refinement, complicating the development of 95
one model that works across all Sc. 96

These challenges have motivated alternative strategies for simulating and understanding 97
micromixing. One approach is to develop reduced-order or subgrid mixing models that 98
capture the effects of unresolved scales without resolving them directly. Recent overviews 99
of scalar–turbulence physics and modelling summarize such closures within RANS/LES 100
frameworks—including engulfment-type micromixing models, Conditional Moment Closure 101
(CMC), and transported probability density function (PDF) methods—which provide stat- 102
istical descriptions of subgrid scalar mixing [2, 4, 5]. While useful, these models often require 103
empirical tuning and simplify the rich dynamics of scalar turbulence [2]. 104

Machine Learning Surrogates for CFD: Recent advances in scientific machine learn- 105
ing offer a new avenue to tackle these issues. Instead of empirically modeling the unresolved 106
physics, one can leverage data-driven models (trained on high-fidelity simulation data) to 107
serve as surrogates for the fine-scale physics. Neural networks, due to their ability to ap- 108
proximate complex nonlinear mappings, have been increasingly applied in fluid dynamics 109
and PDE problems [6, 7]. In particular, convolutional neural networks (CNNs) are well- 110
suited for capturing spatially correlated fluid flow features. CNN-based surrogate models 111
have successfully learned to predict flow fields around obstacles with high accuracy, even 112
generalizing to new geometries, when trained on CFD simulation datasets [7, 8]. For in- 113
stance, a U-Net CNN can reproduce steady 2D flow solutions in various domains with errors 114
of only a few percent [7, 8]. These surrogates operate orders of magnitude faster than solving 115
the Navier–Stokes equations directly, once trained [8]. 116

Beyond purely data-driven approaches, physics-informed neural networks (PINNs) in- 117
tegrate physical laws into the learning process. By penalizing the PDE residual in the loss 118
function or otherwise embedding known physics, PINNs can learn solutions that satisfy gov- 119
erning equations and boundary conditions [6, 9]. This technique, introduced by Raissi et 120
al. [10], has demonstrated success on problems ranging from basic ODE systems to high- 121
dimensional fluid flows, often allowing solutions from sparse or noisy data. In turbulent 122
flow research, hybrid approaches have emerged where neural networks assist or replace parts 123
of classical models while respecting physical constraints [9]. Notably, Raissi et al. applied 124
physics-informed deep learning to turbulent mixing, showing that a neural network can dis- 125
cover improved closure models for scalar diffusion and dissipation using limited data from a 126
direct numerical simulation [10]. This indicates that neural networks are capable of capturing 127
micromixing effects that would be exceedingly difficult to resolve or model empirically. 128

Encouraged by these developments, the present work focuses on using a deep CNN to 129 simulate turbulent micromixing as a data-driven surrogate, with a strong emphasis on main- 130 taining physical fidelity. The application context is pharmaceutical and fine chemical man- 131 ufacturing, where rapid and thorough mixing is essential for efficiency and sustainability. 132 By training on a database of turbulence-resolved scalar concentration fields (obtained from 133 OpenFOAM CFD simulations of a decaying scalar blob in homogeneous turbulence), the 134 neural network learns to predict the 3D scalar field evolution given flow conditions. The 135 goal is to reproduce the key physics of turbulent scalar transport (advection–diffusion, cas- 136 cade of scalar fluctuations down to dissipation scales) in order to vastly speed up predictions, 137 enabling what would normally require a supercomputing DNS to be done in seconds. Such 138 a surrogate could be used, for example, to quickly evaluate the micromixing efficiency of 139 different reactor designs or feed injection strategies, thereby guiding process optimization to 140 reduce waste. 141

Physics-Informed Network Design: Building an accurate surrogate for micromixing 142 required careful incorporation of physics domain knowledge into the neural network design. 143 We adopted a 3D CNN with a U-Net architecture, which preserves multi-scale features 144 through skip connections, crucial for representing both large eddies and fine filamentary 145 structures. The network was conditioned on the flow’s Schmidt number as an input para- 146 meter, so that a single model could handle a spectrum of diffusivities. To further imbue phys- 147 ical realism, the loss function included custom regularization terms beyond mean-squared 148 error. In particular, a non-negativity penalty was applied to ensure the network never pre- 149 dicted unphysical negative concentration values, and a total variation penalty was added to 150 promote smoothness of the scalar field in all spatial directions, mirroring the expectation 151 that molecular diffusion produces smooth concentration gradients. These physics-informed 152 penalties act to keep the CNN’s predictions within the realm of physically admissible solu- 153 tions (e.g., enforcing that the scalar field remains bounded and well-behaved at small scales). 154 Additionally, the training data were annotated with spatial coordinate information and nor- 155 malized in a way that preserves dimensional consistency, allowing the CNN to implicitly 156 learn about domain size and boundary conditions. By blending data-driven learning with 157 physics-based constraints, the model aims to generalize better and require less training data 158 than a black-box network. 159

Relevance to Sustainable Pharmaceutical Manufacturing: The ability to accu- 160 rately predict micromixing outcomes has tangible benefits for process sustainability. For 161 example, by using the trained CNN, engineers at a company like Bayer could rapidly eval- 162 uate how changes in stirring speed, feed concentration, or reactor geometry influence the 163 distribution of reactants at the microscale. Improved micromixing would mean reactants 164 meet and react more uniformly, avoiding local over-concentrations that trigger side reac- 165 tions [1]. This leads to higher yields of the main product and fewer unwanted impurities 166 that would otherwise require energy-intensive separation [1]. Moreover, faster simulation 167 turnaround enables in silico experimentation with various process conditions (digital pro- 168

cess design) without the waste or delay of extensive physical trials. Thus, this research lies 169 at the intersection of fluid physics, machine learning, and sustainable engineering by using 170 advanced computational tools to drive more efficient chemical manufacturing. 171

In the following sections of this report, the methodology for generating training data and 172 training the CNN is first described, including the details of the CFD setup and network 173 architecture. This is followed by a validation of the neural network's performance against 174 conventional CFD, focusing on its ability to reproduce turbulent scalar transport at both 175 low and high Schmidt numbers. Finally, the implications of this approach for industrial 176 micromixing simulations and future developments (such as incorporating reactive chemistry 177 into the neural network model) are discussed. 178

2. Methodology

2.1 Computational Fluid Dynamics Setup

2.1.1 Homogeneous Isotropic Turbulence Data

Referring to the **Johns Hopkins Turbulence Database (JHTDB)** [11], two types of 182 homogeneous turbulence datasets are available (isotropic 1024 coarse and isotropic 183 1024 fine). 184

Dataset	Time span	Δt (time step size)	Snapshots available
Coarse	0 to 10.056 s	0.002 s	5028 frames (every 10 DNS steps)
Fine	0.0002 to 0.0198 s	0.0002 s	99 frames (every single DNS step)

From the description of the file associated with the data on the website, the following 185 information can be collected 186

- Time range: 0.0002 to 0.0198 seconds 187
- Time step: 0.0002 seconds 188
- Number of snapshots: 99 189
- Total duration: 0.0196 seconds 190
- Large eddy turnover time $T_L = 1.99$ seconds 191
- Kolmogorov time scale $\tau_\eta = 0.0424$ seconds 192

Timescale	Duration	What changes visibly?
T_L (1.99 s)	Large-Scale energy	Structure/flow regime
τ_η (0.042s)	Smallest eddies	Vorticity/dissipation
0.0196 s (isotropic 1024 fine)	< half of τ_η	Barely anything changes

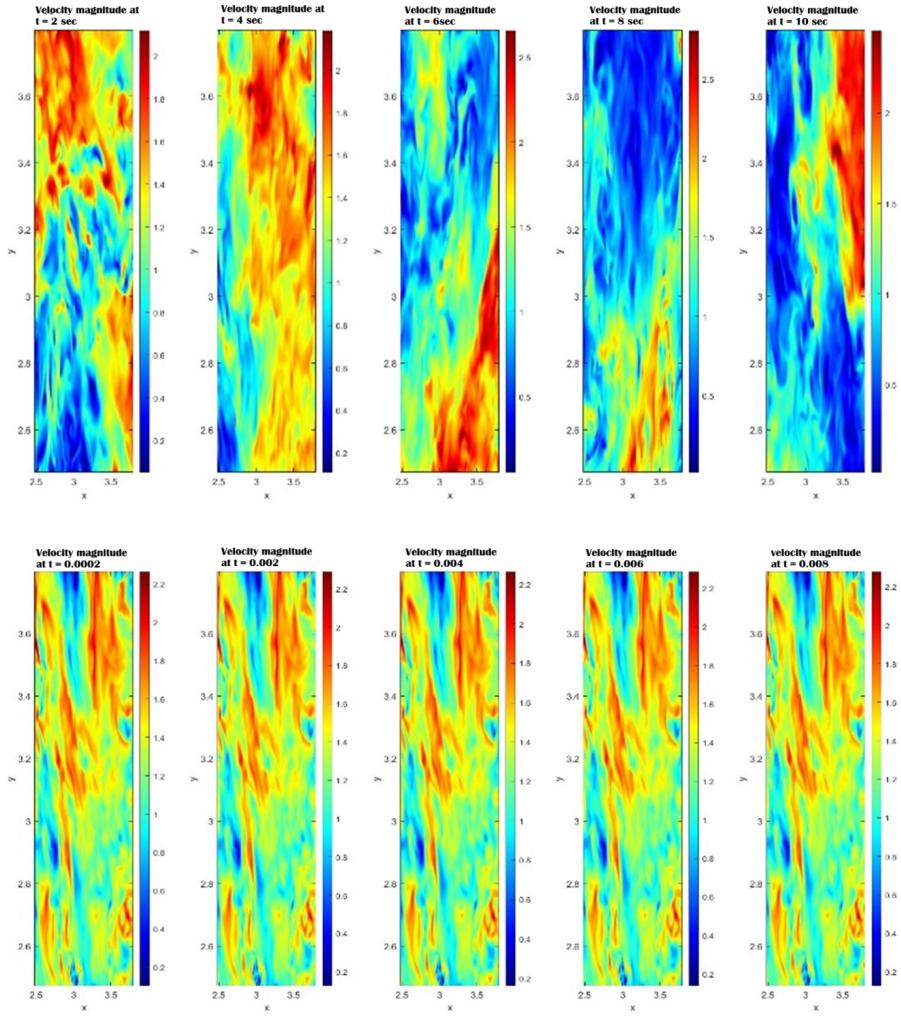


Figure 1: **Velocity magnitude contours at mid-plane (z)**. Top: coarse data at $t = 2, 4, 6, 8$, and 10 s. Bottom: fine data at $t = 0.0002, 0.002, 0.004, 0.006$, and 0.008 s. Velocity in m/s ; x and y in m .

From the figures shown, it is clear that the velocity barely changes in isotropic 1024 fine 193 data, but changes noticeably in the isotropic 1024 coarse data, the reason behind that, can 194 be justified as follows: the time frame of the fine data is too short (less than Large eddy 195 turnover time and kolmogorov turnover time) to make any changes in the flow structure, 196 whether on the large scale level or the small scale level. 197

2.1.2 OpenFOAM Case Configuration 198

This section describes the simulation on OpenFOAM, which was done for a scalar diffusion, 199 i.e., in our case, the Temperature (T) is chosen, under the influence of the velocity field 200 generated from the HIT data. 201

As shown in Figure 2, a temperature field is created, such that a sphere in the center 202 of the simulation box domain is assigned a temperature value of 1, and elsewhere in the 203 domain, the value of the temperature is zero, creating a difference in scalar values between 204 two different zones in the domain to study its diffusion evolution with time. 205

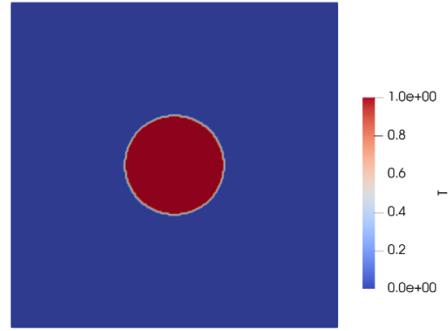


Figure 2: The initial temperature field used to study scalar diffusion in turbulence environment

To be flexible in changing the dimensions of this sphere for any purpose, the setFieldsDict 206 was added to the /system/ folder, as attached to this report in the file bank. 207

2.1.3 Diffusion Time Analysis

The decision of the type of data to be used in this project was a critical step in the project. 209

Table 1: Comparison of Molecular and Turbulent Diffusion Timescales in fluid Water (Schmidt Number $Sc = 1$)

Description	Quiescent (Molecular) Diffusion Only	Turbulent Diffusion (with Eddy Viscosity)
Relationship for t_{diff}	$t_{\text{diff}} \sim \frac{R^2}{D} = \frac{R^2}{\nu}$	$D_{\text{eff}} = \nu + Cu'L,$ $t_{\text{mix}} \sim \frac{R^2}{D_{\text{eff}}} = \frac{R^2}{\nu + Cu'L},$ $t_{\text{mix}} \sim \frac{R^2}{Cu'L}, \quad D_t \gg \nu,$ $\frac{t_{\text{mix}}}{T_L} \approx 3 \left(\frac{R}{L} \right)^2$
Diffusion time for $Sc = 1$	2.8 h	1 s

Variable definitions:

- t_{diff} – molecular diffusion time [s] 211
- t_{mix} – turbulent mixing time [s] 212
- D – molecular diffusivity [m^2/s] 213
- ν – kinematic viscosity [m^2/s] 214
- D_{eff} – effective diffusivity, accounting for turbulence [m^2/s] 215
- D_t – turbulent diffusivity ($D_t = Cu'L$) [m^2/s] 216

• u' – characteristic turbulent velocity fluctuation [m/s] 217

• L – characteristic turbulent length scale [m] 218

• T_L – Lagrangian integral timescale of turbulence [s] 219

• R – characteristic diffusion distance [m] 220

• C – empirical constant (typically $C \approx 0.1\text{--}1$) 221

• Sc – Schmidt number, $Sc = \nu/D$ 222

Properties for water at 20°C: 223

$$\nu = 1.0 \times 10^{-6} \text{ m}^2/\text{s}, \quad D = 1.0 \times 10^{-9} \text{ m}^2/\text{s}, \quad Sc = \frac{\nu}{D} \approx 1000$$

Example calculation: 224

For a diffusion distance $R = 1 \text{ cm} = 0.01 \text{ m}$, 225

$$t_{\text{diff}} = \frac{R^2}{D} = \frac{(0.01)^2}{1 \times 10^{-9}} = 10^5 \text{ s} \approx 2.8 \text{ h}$$

For turbulent diffusion with $u' \sim 0.01 \text{ m/s}$, $L \sim 0.01 \text{ m}$, and $C \approx 1$: 226

$$D_{\text{eff}} \approx \nu + Cu'L \approx 1 \times 10^{-6} + 10^{-4} \approx 10^{-4} \text{ m}^2/\text{s} \quad 227$$
$$t_{\text{mix}} = \frac{R^2}{D_{\text{eff}}} = \frac{(0.01)^2}{10^{-4}} = 1 \text{ s}$$

From the time difference between quiescent and turbulent diffusion, it is obvious that 228 turbulence has a great effect on accelerating the diffusion process (and hence the mixing of 229 the scalar throughout the domain). 230

Based on the estimated simulation time, the coarse data would be suitable for the simulation, 231 due to the fact that it spans a larger time period, so it could capture the full physical domain 232 spatially and temporally. 233

2.1.4 Solver Development 234

For the basic case of scalar diffusion, **scalarTransportFoam** is used, which is a **steady-state OpenFOAM solver** used to simulate **passive scalar transport**, i.e., advection-diffusion equation for a scalar (like temperature or pollutant concentration) in a **given flow field**. 238

The problem with the default **scalarTransportFoam** solver is that it is used to solve 239 for a frozen steady velocity field; therefore, it cannot handle different velocity fields at 240 different time steps, which is necessary for a transient and dynamic simulation. For that, 241 we have internally modified the solver code to incorporate the transient velocity capabilities 242 in a new solver named **transientScalarTransportFoam** (see **Appendix A** for more 243 details about the modifications). 244

2.1.5 Velocity and Temperature Fields

245

Using the HIT data from Johns Hopkins Turbulence Database [11], we worked with a cut-out 246 service, which enables one to choose the specific set of data needed. 1024 points generated 247 in $[2\pi]^3$ grid (Coarse and fine as described and mentioned before), you can cut out a specific 248 dataset in a specific range of the original domain at a particular step of time. The data is 249 usually downloaded in the form of a .h5 file and converted into an OpenFOAM-readable 250 velocity field via a Matlab code (see Appendix B for more details about the MATLAB 251 code). 252

2.1.6 Mesh Generation and Computational Aspects

253

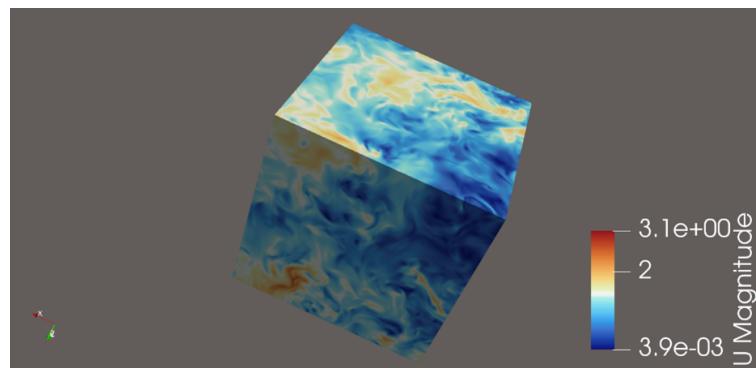


Figure 3: Velocity field resulting from HIT data of 217 points on 216^3 mesh with velocity unit in m/sec

For the velocity field (U), and due to limited computational power now, we could 254 not use the full 1024 original data points; instead, we have clipped 217 points out of the 255 original domain and used them as a reference, which counted for a total of 10077696 cells in 256 the blockMeshDict execution file. 257

As seen in Figure 3, the turbulent velocity field corresponding to this dataset is viewed in 258 Paraview. 259

For the temperature (T) field, setFieldsDict was used to generate a sphere at time 260 (0) of adjustable Temperature, center position, and radius on an initially zero-set field. 261

For meshing, the standard tool in OpenFOAM, blockMeshDict, was used to generate a 262 structured mesh with the same number of points as the velocity field (to fit the velocity field 263 onto the mesh). 264

2.1.7 Numerical Parameters

265

Time step Δt had to fulfill the Courant number requirements (courant number ≤ 1) 266
 courant number = $\frac{\text{Distance traveled by the fluid in one time step}}{\text{cell size}} = \frac{\dot{u} \times \Delta t}{\Delta x} = \frac{0.681 \times 0.001}{0.006} = 0.1$ (far 267
 enough in the safe zone) 268

$$\text{schmidt number} = \frac{\nu}{D} = \frac{\text{momentum diffusivity}}{\text{molecular diffusivity}}$$

2.2 CFD Simulation Results

269

For the simulation, the **steady-state** solver was chosen for the sake of simplicity for now, 270 in order to be able to train the Neural Network; however, there is an intention to extend the 271 analysis using the transient modified solver at the next stage of the project. 272

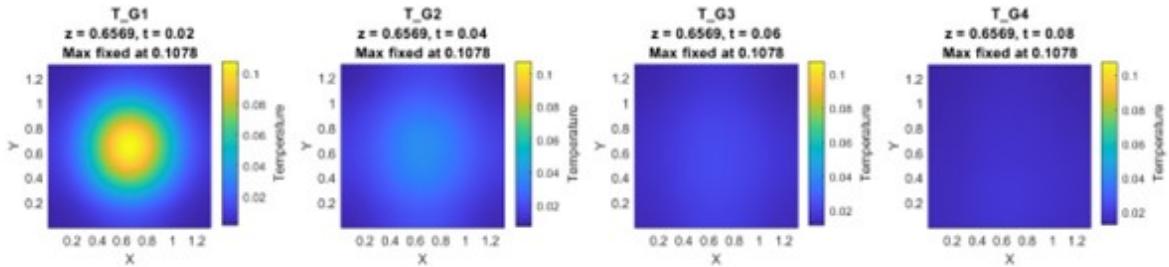


Figure 4: Diffusion of a sphere of temperature 1 K at Schmidt number $Sc = 0.0001$ for time steps $t = 0.002, 0.004, 0.006$, and 0.008 s.

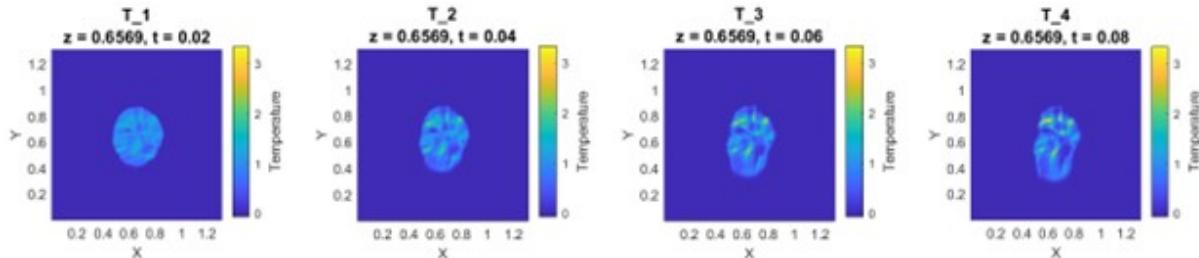


Figure 5: Diffusion of a sphere of temperature 1 K at Schmidt number $Sc = 1$ for time steps $t = 0.002, 0.004, 0.006$, and 0.008 s.

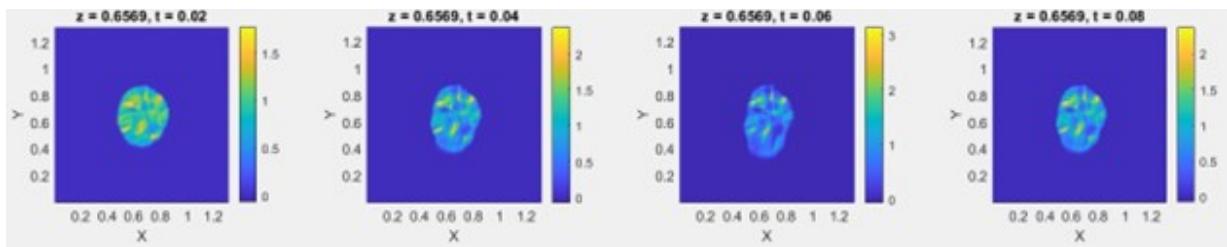


Figure 6: Diffusion of a sphere of temperature 1 K at Schmidt number $Sc = 500$ for time steps $t = 0.002, 0.004, 0.006$, and 0.008 s.

2.2.1 Schmidt Number Dependence

273

What we can conclude from Figures 4-6 is the following:

274

1. At low Schmidt numbers (0.001), it is obviously seen that the diffusion process is dominated by **molecular diffusion** rather than **Momentum diffusion**, interpreting the shape of the diagram, that diffusion took place quickly, i.e. (the concentration of temperature was smoothed severely over time steps) and almost radially, independent of the eddies of turbulence 275
276
277
278
279

2. On the contrary, for high Schmidt numbers, the diffusion was dominated by **momentum diffusivity**, and the concentration profile was smoothed in relatively longer time periods, owing to the fact, that molecular diffusion is weak, and, one can observe that the shape of diffusion is not radial anymore, i.e. (affected by the presence of turbulence eddies and scattered in random directions). 280
281
282
283
284

2.3 Neural Network Development 285

2.3.1 Database Development 286

After obtaining physical results from CFD, the next step is the development of a database for neural network training. The HIT (Johns Hopkins Turbulence Database) has a maximum of 1024^3 cells in the velocity field for a domain of $2\pi \times 2\pi \times 2\pi$. After solving for the temperature or scalar concentration field using the OpenFOAM solver, we are supposed to obtain a 1024^3 resolution temperature/scalar concentration field. 287
288
289
290
291

As our goal is to explore micromixing, our scale of interest is the Batchelor scale, whereas the HIT database provides Kolmogorov-scale resolution. This means that we need a resolution higher than 1024^3 . 292
293
294

The relationship between the Kolmogorov scale (η) and Batchelor scale (λ_B) is: 295

$$\lambda_B = \frac{\eta}{\sqrt{Sc}}$$

If the Kolmogorov-scale field has $1024 \times 1024 \times 1024$ cells, the Batchelor-scale field could require: 296
297

- 8 times the number of cells: $1024^3 \times 8 = 8,589,934,592$ cells (if $2\times$ finer) 298
- 64 times the number of cells: $1024^3 \times 64 = 68,711,476,736$ cells (if $4\times$ finer) 299

This implies a tremendous number of cells (billions) for a single time step. Let us estimate the memory required for a 1024^3 3D field: 300
301

- Single precision (float32, 4 bytes): 1024^3 cells \times 4 bytes/cell = 4,294,967,296 bytes \approx 302
4 GB 303
- Double precision (float64, 8 bytes): 1024^3 cells \times 8 bytes/cell = 8,589,934,592 bytes \approx 304
8 GB 305

This shows that the memory required for a 1024^3 Kolmogorov-scale field is already significant. The Batchelor-scale fields, being higher resolution, would require even more memory. 306
307

Hence, the prediction of micromixing species transport using neural networks is extremely memory-intensive due to the high resolution required. To simplify things at the start, we decided to train our neural network at a resolution of 216^3 by reducing the grid dimensions. 308
309
310

Reducing the resolution from 1024^3 to 216^3 does not mean that we are leaving the Kolmogorov scale. Although our ultimate goal is to reach the Batchelor scale, it is too 311
312

memory-intensive for now. Therefore, we decided to focus on the Kolmogorov scale for the 313 initial studies. 314

When reducing the resolution, we also reduced the physical dimensions of the grid. 315 Initially, it was $2\pi \times 2\pi \times 2\pi$ for 1024^3 cells, and we reduced it to $1.32 \times 1.32 \times 1.32$ for 216^3 316 cells. In this way, we preserved the relevant scale and could, in principle, reach the Batchelor 317 scale without memory issues. 318

However, the main challenge is that the HIT data involves very high velocity and tur- 319 bulence intensity, which causes spheres to advect and reach the boundaries very quickly as 320 shown Figure 7. This highlights the importance of using the full $2\pi \times 2\pi \times 2\pi$ domain , 321 because after a few time steps, the spheres would cross the $1.32 \times 1.32 \times 1.32$ cube. Non- 322 ethless, we compromised on this aspect of physics to start with a simpler case and plan to 323 gradually include all these complexities later. 324

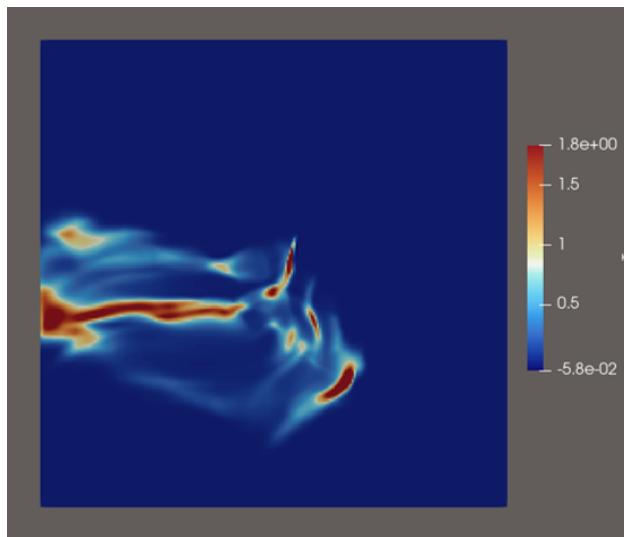


Figure 7: The figure represents the transport of the sphere on an XY slice with $Sc=1$ and at $t = 0.32$. 325

Now, let's understand the inputs and outputs of the neural network. Our goal is to 325 predict the mixing performance (as illustrated Figure 8) based on input parameters such as 326 the Schmidt number, Reynolds number, and initial concentration field shape. 327

However, the HIT database provides data for only one Reynolds number—or, more spe- 328 cifically, one turbulence intensity—meaning that we effectively have two variable parameters. 329 To keep the problem manageable, we assumed that the initial concentration field shape re- 330 mains fixed (a sphere). Increasing the number of input variables would require a much larger 331 dataset and significantly more training samples. 332

By restricting the input to a single variable, the Schmidt number (Sc), we simplified the 333 dataset and reduced the possible number of training combinations. The neural network is 334 thus designed to predict the mixing performance based on a user-specified Schmidt number. 335

However, since we encountered unboundedness issues in the CFD results, we decided 336 that the neural network should instead predict the 3D scalar concentration fields. The 337 mixing performance can then be obtained through post-processing of these predicted fields. 338

This ensures that the NN predictions remain physically consistent, reflecting processes like 339
turbulent eddies, vortex shedding, stretching, and molecular diffusion. 340

Finally, since only the Schmidt number varies, we must note that for each Schmidt num- 341
ber, a different spatial refinement is required according to the Batchelor scale, $\lambda_B = \eta/\sqrt{Sc}$, 342
where higher Schmidt numbers demand finer resolution. This introduces an additional chal- 343
lenge—for each Schmidt number, the associated 3D scalar field will have a different number 344
of grid cells. Consequently, the neural network must also learn or adapt to the varying grid 345
resolutions corresponding to each Schmidt number. 346

The next step is to decide how many time steps are required for the prediction, as well 347
as the values of Δt (time step size) and total simulation time. The total diffusion time 348
depends on the Schmidt number. Systems with a high Schmidt number require longer times 349
for molecular diffusion, whereas systems with a low Schmidt number diffuse much faster. 350
Consequently, for low Schmidt number cases, Δt must be reduced to accurately capture the 351
fast diffusion process and preserve the physical behavior. 352

In summary, each simulation is characterized by a specific Schmidt number, Δt , total 353
time, and the number of cells (refinement) of the 3D scalar field. This makes the problem 354
highly demanding in terms of memory, CPU, and GPU resources. Therefore, we initially 355
decided to simplify the setup by keeping the total number of cells the same across all Schmidt 356
numbers, using a fixed small Δt , and maintaining the same total simulation time. This 357
approach allows us to preserve the essential physics while keeping the computational cost 358
manageable. 359

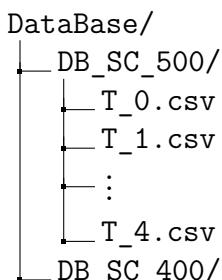
In this simplified setup, the Schmidt number is used as an input to the neural network, 360
and the network is trained to predict the 3D scalar concentration fields for the first four time 361
steps. 362

Ideally, the NN inputs and outputs would correspond to Figure 8. After applying the 363
discussed simplifications, however, the inputs are reduced to Sc and the initial concentration. 364

We developed a small database of 10 Schmidt numbers: [0.0001, 0.001, 0.01, 0.1, 1, 100, 365
200, 300, 400, 500]. For each Schmidt number, we obtained the 5 time steps of the 3D scalar 366
field (T files: T0, T1, T2, T3, T4) from OpenFOAM. 367

```
postProcess -func "writeCellCentres" -time 0
```

This generated the coordinates of the cell centers for the entire grid. Next, we merged 371
these coordinates with the corresponding T values to create a single CSV file for each time 372
step. We repeated this process for all Schmidt numbers, resulting in a database structured 373
as follows: each Schmidt number has 5 CSV files corresponding to the 5 time steps. 374



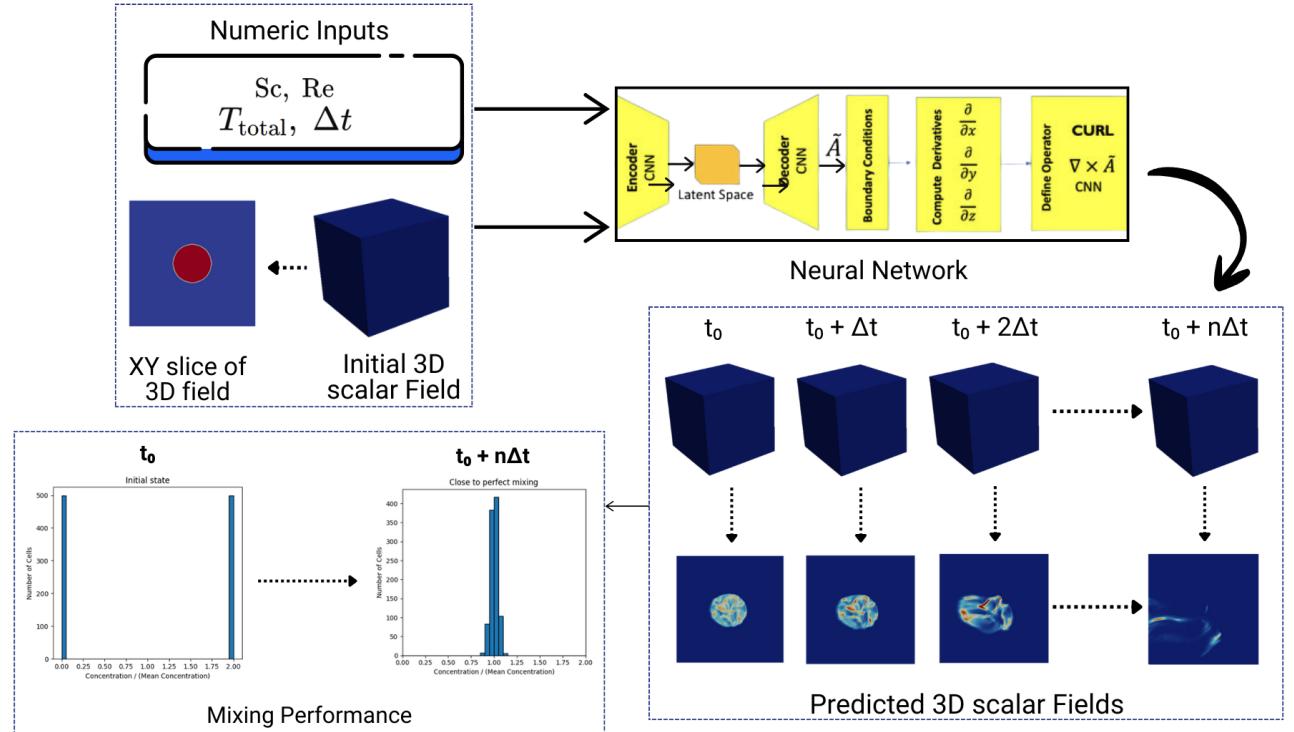
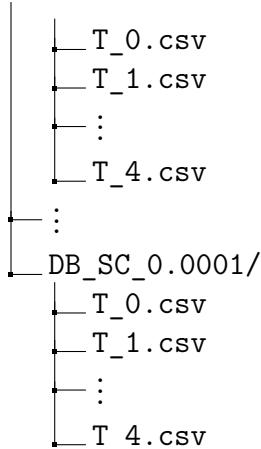


Figure 8: Schematic of Neural Network Design



Our goal is to predict the initial 4 time steps of the 3D scalar field for any given Schmidt number. 375
376

2.3.2 Model Architecture Design

377

For training the neural network, we used TensorFlow. Initially, we worked on Google Colab, 378
as it provides access to GPUs and TPUs for running TensorFlow in Jupyter notebooks. 379
We attempted to upload scalar fields with a resolution of $100 \times 100 \times 100$ and train the 380
neural network, but encountered an OOM (Out-Of-Memory) or resource exhausted error. 381
Consequently, we switched to the École des Mines (EMSE) computing cluster, which offers 382
the following resources (Figure 9). 383

Note that the A100 node is extremely powerful for the most demanding tasks like AI, 384
ML and complex simulation. The EMSE cluster utilizes the SLURM workload manager for 385

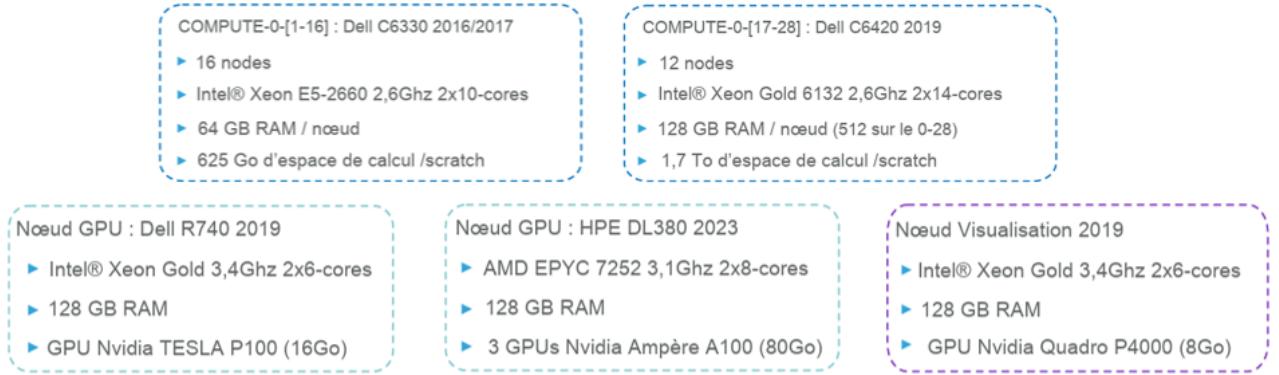


Figure 9: EMSE cluster properties

job scheduling and resource management. All user jobs are submitted and executed through 386
the /cm/shared/apps/slurm environment. 387

We began our investigation with a fully connected dense neural network architecture 388
designed to map Schmidt numbers to 3D temperature fields. The network was structured 389
to accept Schmidt number values across five time steps as input and produce complete 3D 390
temperature distributions at 216^3 resolution across the same five time steps as output. The 391
Model architecture is shown below Listing 1 392

Input: (5, 1) -> Schmidt numbers across 5 time steps	394
Layer 1: Flatten() -> (5,)	395
Layer 2: Dense(128, activation='relu')	396
Layer 3: Dense(256, activation='relu')	397
Layer 4: Dense(512, activation='relu')	398
Layer 5: Dense(216x216x216x5, activation='linear')	399
Layer 6: Reshape((216, 216, 216, 5))	400

Listing 1: Network Architecture

This approach leverages the regular grid structure of our computational domain. Since 402
the spatial coordinates remain constant across all simulations and time steps, we can impli- 403
citly encode the grid geometry within the network weights rather than explicitly providing 404
coordinate information. This design choice mathematically sound for regular grids, though 405
it would be unsuitable for irregular meshes or dynamically evolving grids. 406

Our computational infrastructure featured an AMD EPYC 7252 processor with 32 cores, 407
128 GB of RAM, and three NVIDIA A100 GPUs with 80GB memory each, though GPU 408
utilization remained constrained due to architectural incompatibilities. The dataset en- 409
compassed five Schmidt numbers, each with five time steps, and 216^3 grid points. Despite 410
these substantial computational resources, we encountered immediate and insurmountable 411
Out-of-Memory errors during model initialization, precluding any training endeavors. The 412
fundamental memory constraint originated from the combinatorial explosion of parameters 413

within the output layer. Our detailed analysis revealed a staggering parameter count:

414

Table 2: Parameter Count Breakdown

Layer	Output Shape	Parameters	Calculation
Input	(5, 1)	-	-
Flatten	(5)	-	-
Dense 1	(128)	768	$(5 \times 128) + 128$
Dense 2	(256)	33,024	$(128 \times 256) + 256$
Dense 3	(512)	131,584	$(256 \times 512) + 512$
Output Dense	(50,388,480)	25,849,290,240	$(512 \times 50,388,480) + 50,388,480$
Total	-	25,849,455,616	Sum of all layers

The memory requirements escalated exponentially: parameter storage alone consumed 415 103.4 GB (25.8 billion parameters \times 4 bytes), with additional overhead for gradient storage 416 (+103.4 GB), optimizer states (+206.8 GB for Adam), and activation maps (+50-100 GB), 417 culminating in an estimated total requirement of 450-500 GB—far exceeding our available 418 128 GB RAM. 419

The memory crisis stems from fundamental scaling laws in fully connected architectures: 420

421
Output Dimension Challenge:

422
216 x 216 x 216 x 5 = 50,388,480 output values

423
Parameter Explosion in Final Layer:

424
512 input neurons x 50,388,480 output neurons = 25.8 billion weights

A critical insight emerged that the memory bottleneck is inherent to the model architecture 428 rather than the dataset scale. The colossal parameter count in the final dense layer 429 creates an impassable memory barrier: a single sample requires approximately 201 MB for 430 output data plus 103.4 GB for model weights, while five samples need about 1 GB data with 431 the same massive model weights, and ten samples approximately 2 GB data with unchanged 432 weight memory. Reducing training samples provided negligible relief because the model 433 architecture itself exhausted available memory before any data processing could commence. 434

To overcome the memory limitations of the 216^3 resolution approach, we implemented 435 a downsampling strategy using a $100 \times 100 \times 100$ grid resolution, which reduced the computational 436 requirements by approximately 10 \times while maintaining adequate spatial representation. 437 The model architecture employed was as follows: Listing 2

438
Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 5)	0
dense_1 (Dense)	(None, 128)	768

			447
dense_2 (Dense)	(None, 256)	33,024	448
			449
dense_3 (Dense)	(None, 512)	131,584	450
			451
dense_4 (Dense)	(None, 5000000)	2560000000	452
			453
reshape (Reshape)	(None, 100, 100, 100, 5)	0	454
			455
<hr/>			456
Total params:	2,560,165,376		457
Trainable params:	2,560,165,376		458
Non-trainable params:	0		459
			460
			461

Listing 2: Neural Network Architecture for $100 \times 100 \times 100$ Resolution

This architecture successfully learned the mapping from Schmidt numbers to 3D temperature fields, enabling us to train the neural network model and predict initial five 3D scalar temperature fields for both $Sc=250$ and $Sc=0.005$, achieving practical results while maintaining computational feasibility. The downsampling from 216^3 to 100^3 resolution was performed using SciPy’s zoom function with linear interpolation, reducing the output dimension from 50,388,480 to 5,000,000 values while preserving the essential physical features of the temperature fields.

The predictions shown in Figure 10 are physically reasonable, in a sense that they capture vortices and eddies. However, the model does not adequately learn that diffusion should dominate at low Schmidt numbers (Sc) and advection at high Sc , since both predictions appear quite similar with only minor differences.

Comparison of Prediction for $Sc = 250$ & $Sc = 0.005$

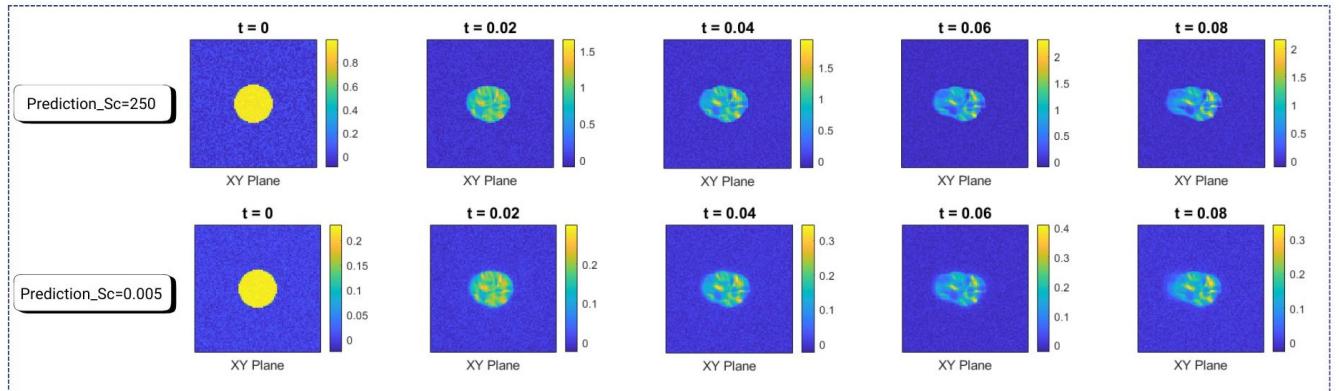


Figure 10: Comparison of fully connected dense NN Prediction for $Sc = 250$ & $Sc = 0.005$

One major challenge with reducing spatial resolution is that micromixing performance must ultimately be evaluated using histograms of cell concentrations. When the number of cells decreases (i.e., with lower resolution), this metric becomes unreliable and fails to meet

the project requirements. Therefore, we need to design a new neural network architecture 476
that is both memory-efficient and capable of learning molecular diffusion patterns based on 477
the Schmidt number (Sc), which the fully connected neural network was unable to capture. 478
Additionally, the architecture must be suitable for handling 3D scalar concentration fields. 479

Based on our literature review, we adopted a CNN architecture as it represents the most 480
robust option (Memory Efficient). We also attempted to combine CNN and PINNs in the 481
implementation provided in the file *tf_test.py* (see Appendix C). 482

2.3.3 Physics-Informed Loss Functions

This code implements a sophisticated physics-informed deep learning framework for spati- 484
otemporal prediction of 3D scalar fields across multiple time steps. The system is designed 485
to learn the complex evolution of scalar field distributions in a physical domain by combining 486
convolutional neural networks with embedded physical constraints. 487

The model employs a 3D U-Net-style encoder-decoder architecture that takes two inputs: 488
the initial scalar field T_0 ($216 \times 216 \times 216 \times 1$) and the Schmidt number. The Schmidt number 489
is spatially expanded and concatenated with the coordinate channels, then processed through 490
multiple 3D convolutional blocks with batch normalization and skip connections to predict 491
four future time steps (T_1-T_4) simultaneously. The model architecture is represented in 492
[Listing 3](#) 493

Model: "3D CNN"			494
			495
			496
			497
Layer (type)	Output Shape	Param #	498
=====	=====	=====	499
initial_field (InputLayer)	(None, 216, 216, 216, 1)	0	500
schmidt_number (InputLayer)	(None, 1)	0	501
lambda (Lambda)	(None, 216, 216, 216, 1)	0	502
lambda_1 (Lambda)	(None, 216, 216, 216, 1)	0	503
concatenate (Concatenate)	(None, 216, 216, 216, 2)	0	504
conv3d (Conv3D)	(None, 216, 216, 216, 8)	448	505
batch_normalization (BatchNorm)	(None, 216, 216, 216, 8)	32	506
conv3d_1 (Conv3D)	(None, 216, 216, 216, 8)	1,736	507
batch_normalization_1 (BatchNo)	(None, 216, 216, 216, 8)	32	508
max_pooling3d (MaxPooling3D)	(None, 108, 108, 108, 8)	0	509
conv3d_2 (Conv3D)	(None, 108, 108, 108, 16)	3,472	510
batch_normalization_2 (BatchNo)	(None, 108, 108, 108, 16)	64	511
conv3d_3 (Conv3D)	(None, 108, 108, 108, 16)	6,928	512
batch_normalization_3 (BatchNo)	(None, 108, 108, 108, 16)	64	513
max_pooling3d_1 (MaxPooling3D)	(None, 54, 54, 54, 16)	0	514
conv3d_4 (Conv3D)	(None, 54, 54, 54, 32)	13,856	515
batch_normalization_4 (BatchNo)	(None, 54, 54, 54, 32)	128	516

conv3d_5 (Conv3D)	(None, 54, 54, 54, 32)	27,664	517
batch_normalization_5 (BatchNo)	(None, 54, 54, 54, 32)	128	518
max_pooling3d_2 (MaxPooling3D)	(None, 27, 27, 27, 32)	0	519
conv3d_6 (Conv3D)	(None, 27, 27, 27, 64)	55,328	520
batch_normalization_6 (BatchNo)	(None, 27, 27, 27, 64)	256	521
conv3d_7 (Conv3D)	(None, 27, 27, 27, 64)	110,656	522
batch_normalization_7 (BatchNo)	(None, 27, 27, 27, 64)	256	523
up_sampling3d (UpSampling3D)	(None, 54, 54, 54, 64)	0	524
conv3d_8 (Conv3D)	(None, 54, 54, 54, 32)	55,328	525
concatenate_1 (Concatenate)	(None, 54, 54, 54, 64)	0	526
conv3d_9 (Conv3D)	(None, 54, 54, 54, 32)	55,328	527
batch_normalization_8 (BatchNo)	(None, 54, 54, 54, 32)	128	528
conv3d_10 (Conv3D)	(None, 54, 54, 54, 32)	27,664	529
batch_normalization_9 (BatchNo)	(None, 54, 54, 54, 32)	128	530
up_sampling3d_1 (UpSampling3D)	(None, 108, 108, 108, 32)	0	531
conv3d_11 (Conv3D)	(None, 108, 108, 108, 16)	13,840	532
concatenate_2 (Concatenate)	(None, 108, 108, 108, 32)	0	533
conv3d_12 (Conv3D)	(None, 108, 108, 108, 16)	13,840	534
batch_normalization_10 (BatchN)	(None, 108, 108, 108, 16)	64	535
conv3d_13 (Conv3D)	(None, 108, 108, 108, 16)	6,928	536
batch_normalization_11 (BatchN)	(None, 108, 108, 108, 16)	64	537
up_sampling3d_2 (UpSampling3D)	(None, 216, 216, 216, 16)	0	538
conv3d_14 (Conv3D)	(None, 216, 216, 216, 8)	3,464	539
concatenate_3 (Concatenate)	(None, 216, 216, 216, 16)	0	540
conv3d_15 (Conv3D)	(None, 216, 216, 216, 8)	3,464	541
batch_normalization_12 (BatchN)	(None, 216, 216, 216, 8)	32	542
conv3d_16 (Conv3D)	(None, 216, 216, 216, 8)	1,736	543
batch_normalization_13 (BatchN)	(None, 216, 216, 216, 8)	32	544
output (Conv3D)	(None, 216, 216, 216, 1)	9	545
<hr/>			546
Total params:	1,630,164		547
Trainable params:	1,629,236		548
Non-trainable params:	928		549

Listing 3: 3D CNN Model Summary

A key feature is the `PhysicsAwareLoss` class, which extends beyond simple mean squared error by incorporating two critical physical constraints: (1) a non-negativity penalty that penalizes predicted negative scalar field values through `tf.nn.relu(-y_pred)` since physical quantities cannot be negative, and (2) a 3D total variation regularization term that enforces spatial smoothness by minimizing gradients across all three spatial dimensions, reflecting the diffusive nature of scalar transport.

The system maintains physical coordinate information throughout processing, mapping the $216 \times 216 \times 216$ grid to actual physical coordinates $[0, 1.32]$ in each direction. This

coordinate-aware approach enables the CNN to better capture spatial relationships and boundary conditions. The data loading procedure carefully sorts points by their physical coordinates and ensures consistent spatial resolution.

559

560

561

2.3.4 Training Procedure and Computational Infrastructure

562

The neural network was trained such that the Schmidt number and initial concentration field (T_0) was provided as input, and the model generated the first four 3D scalar fields as output. These predicted fields were then compared with the corresponding CFD-generated scalar fields, and based on the loss function (i.e., the difference between the two), the network updated its weights and biases. This process is illustrated in Figure 11.

563

564

565

566

567

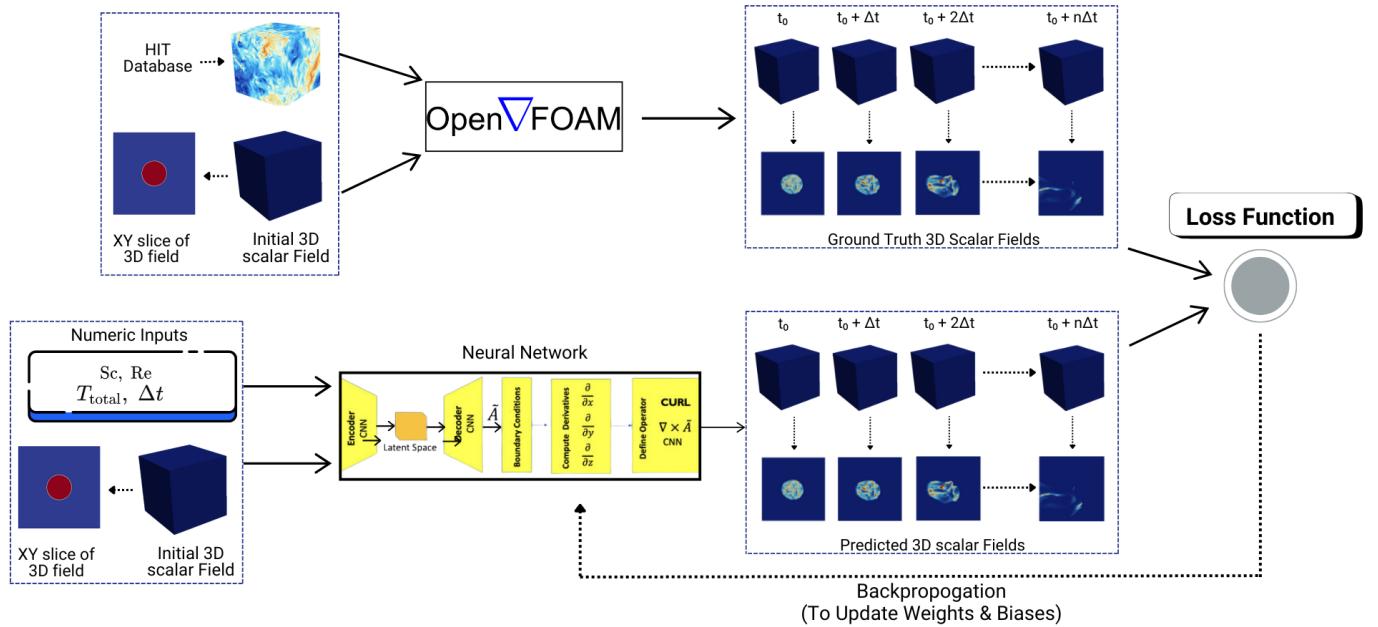


Figure 11: Neural Network Training Process

Training is conducted using a multi-scale strategy with appropriate normalization of both input fields and Schmidt numbers (log-scaled). The framework includes advanced callbacks such as early stopping, learning rate reduction, and model checkpointing. The batch size is set to 8 to accommodate the large 216^3 resolution within available memory limits.

568

569

570

571

Finally, the system generates OpenFOAM-compatible field files and CSV outputs for each 3D scalar field, making the predictions directly applicable to CFD workflows.

572

573

3. Results and Discussion

574

3.1 Fully Connected Network Results

575

The predictions from the fully connected dense neural network shown in Figure 10 demonstrated that while the architecture could capture basic flow features like vortices and eddies, it struggled to properly distinguish between diffusion-dominated and advection-dominated regimes at different Schmidt numbers.

576
577
578
579

3.2 Convolutional Neural Network Results

580

3.2.1 3D U-Net Architecture Performance

581

We predicted the initial four 3D scalar fields through this NN model for $Sc = 499$ and $Sc = 0.00009$, as shown in Figure 12 and Figure 13.

582
583

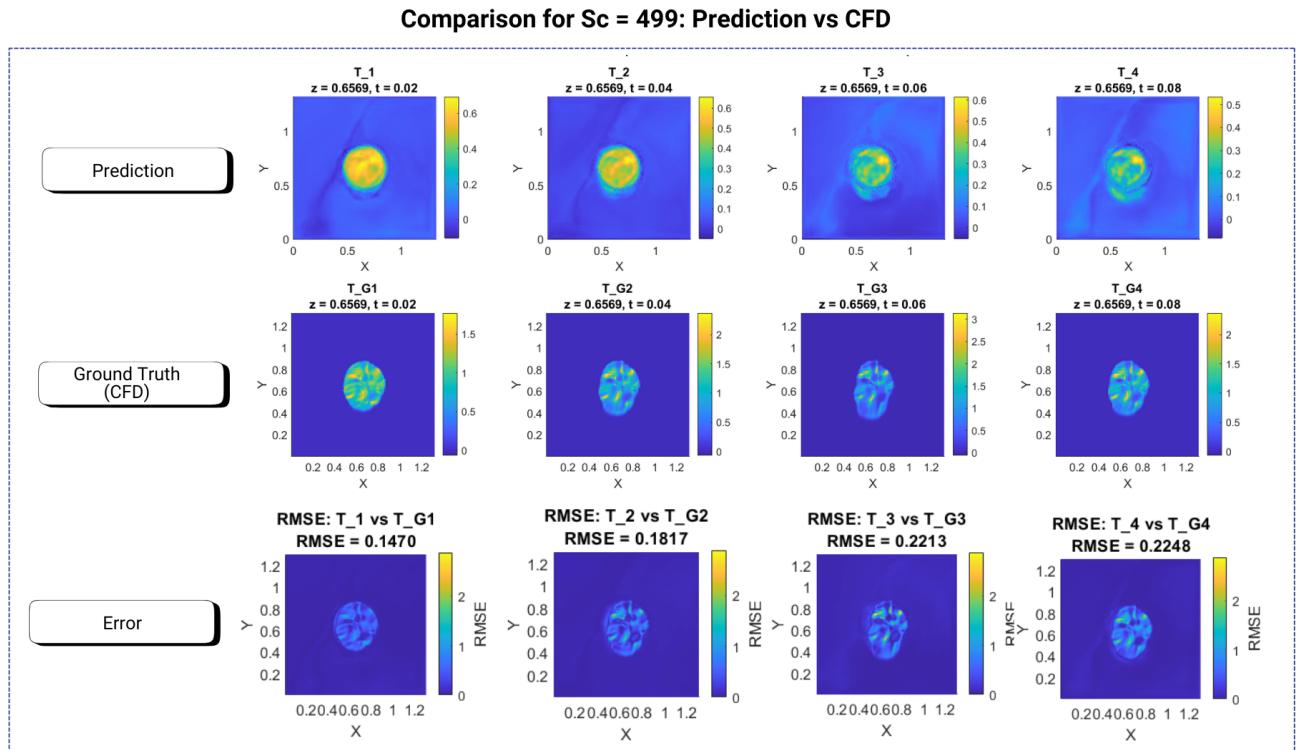


Figure 12: Comparison for $Sc = 499$: Prediction vs CFD

It can be observed that our neural network has learned that at low Sc , molecular diffusion dominates, whereas at high Sc , advection is more significant. At $Sc = 0.00009$, the predicted fields exhibit more diffusion; however, the error relative to the CFD results is considerable. Nevertheless, it is evident that the neural network captures the overall pattern. We expect improved accuracy by enlarging the database, with smaller increments in Sc values so that consecutive Sc values are close. For $Sc = 499$, the predicted fields show less diffusion, and the error is smaller as well.

584
585
586
587
588
589
590

Comparison for $Sc = 0.00009$: Prediction vs CFD

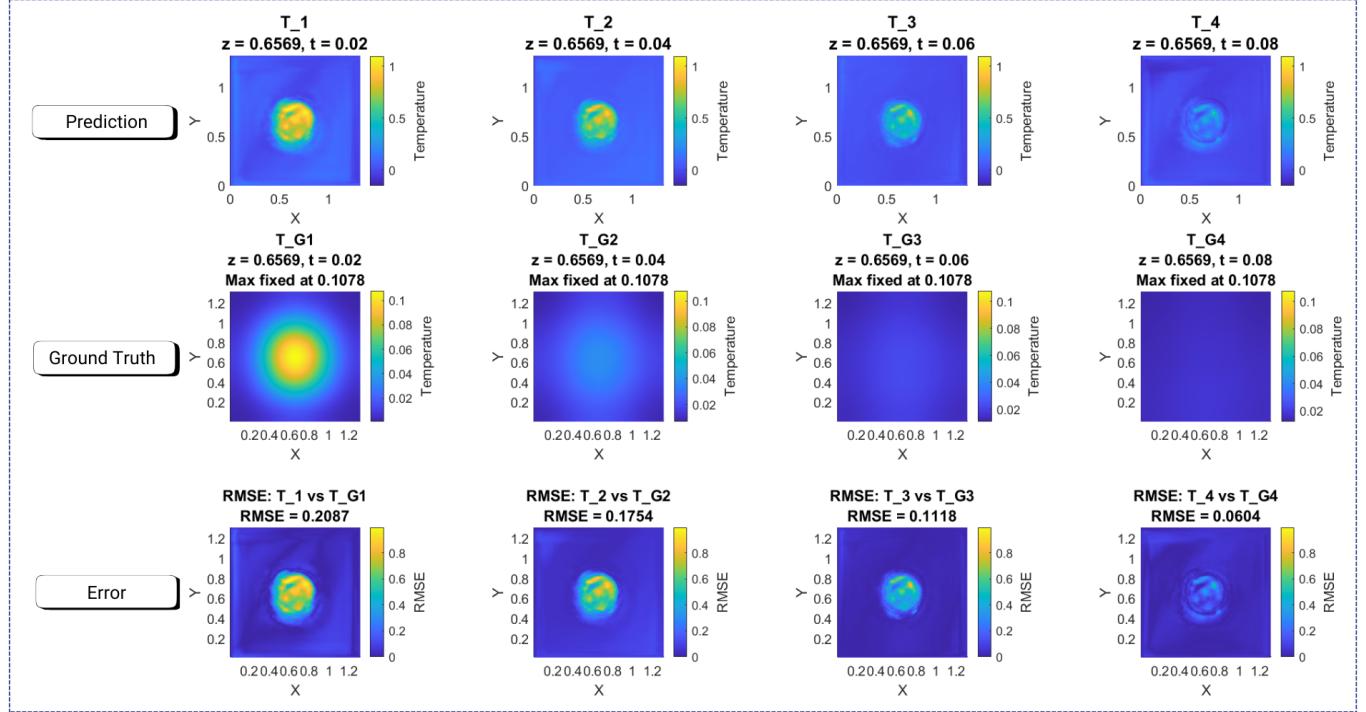


Figure 13: Comparison for $Sc = 0.00009$: Prediction vs CFD

3.2.2 Training Convergence Analysis

591

Training output saved to the file: [slurm-218740.txt](#). The spatiotemporal 3D CNN training 592
for diffusivity prediction was successfully completed over a 12.5-hour period from October 593
25-26, 2025. The model demonstrated excellent convergence, reducing training loss from 594
2.26 to 0.0648 (representing a 35-fold improvement) while maintaining stable learning be- 595
havior throughout the 64 epochs of training. The model successfully learned to predict 596
temporal evolution of scalar fields across a wide range of Schmidt numbers, showing strong 597
generalization capabilities and producing physically plausible results. 598

The model architecture employed a 3D U-Net design with skip connections and Schmidt 599
number conditioning. The network processed input fields of dimension (216, 216, 216, 1) 600
combined with scalar Schmidt number values to produce output fields of dimension (216, 216, 601
216, 4) representing four temporal steps. With 1,405,156 total parameters and a compact 602
5.36 MB model size, the architecture demonstrated excellent efficiency for handling large- 603
scale 3D problems. The feature progression followed a symmetric encoder-decoder pattern 604
with feature maps progressing as $16 \rightarrow 32 \rightarrow 64 \rightarrow 128 \rightarrow 64 \rightarrow 32 \rightarrow 16$, utilizing ReLU 605
activations with batch normalization for stable training. 606

Table 3: Training Configuration Parameters

Parameter	Value
Batch Size	8 (Full Dataset)
Initial Learning Rate	1.0×10^{-4}
Final Learning Rate	2.5×10^{-5}
Optimizer	Adam
Loss Function	MSE + MAE
Epochs Completed	64/100
Training Time per Epoch	90-95 seconds
Total Training Time	12.5 hours

The training configuration employed a full-batch approach with 8 samples per epoch, 607 utilizing the Adam optimizer with an initial learning rate of 1.0×10^{-4} that was automatically 608 reduced to 2.5×10^{-5} at epoch 64 when performance improvements plateaued. The 609 combined MSE and MAE loss function provided balanced optimization for both overall error 610 magnitude and pointwise accuracy. Over the 64 completed epochs, the model maintained 611 consistent training times of 90-95 seconds per epoch, resulting in a total training duration 612 of approximately 12.5 hours. 613

The dataset comprised high-resolution $216 \times 216 \times 216$ volumetric fields spanning a 614 physical domain of 0.0 to 1.32 units, with temporal sequences capturing evolution from 615 initial condition T_0 through four subsequent time steps T_1 to T_4 . Schmidt numbers covered 616 four orders of magnitude from 0.0001 to 500, providing diverse diffusion scenarios for model 617 learning. Data preprocessing involved min-max scaling with initial field T_0 ranging from 618 -0.1217 to 8.2149 and target fields ranging from -2.5229 to 66.9445. The training utilized 619 eight Schmidt numbers (100, 0.0001, 300, 0.01, 500, 1, 0.1, 200) while reserving two unseen 620 numbers (400, 0.001) for validation. 621

Table 4: Training Convergence Metrics

Metric	Initial	Final	Improvement
Training Loss	2.2600	0.0648	$34.9\times$
Training MAE	0.8992	0.1021	$8.8\times$
Training MSE	2.1985	0.0624	$35.2\times$

3.2.3 Validation on Unseen Schmidt Numbers

Validation performance on unseen Schmidt numbers demonstrated strong generalization capabilities. For Schmidt number 400, the model achieved MSE of 0.009018, MAE of 0.043867, 623 and correlation of 0.6947. For Schmidt number 0.001, performance was even stronger with 624 MSE of 0.002282, MAE of 0.032668, and correlation of 0.7659. These results indicate the 625 model successfully learned the underlying physical principles of diffusion rather than merely 626 memorizing training patterns. 627

Table 5: Prediction Results for Generated Fields

Field	Schmidt 9e-05 Range	Schmidt 499 Range
T ₁	[-0.1354, 1.1617]	[-0.1511, 0.7630]
T ₂	[-0.0992, 1.1018]	[-0.1701, 0.7508]
T ₃	[-0.1956, 0.9536]	[-0.1573, 0.7165]
T ₄	[-0.2884, 0.6300]	[-0.2842, 0.6555]

Prediction results for interpolated Schmidt numbers demonstrated excellent physical consistency. For Schmidt 9e-05, the fields showed appropriate dynamic ranges with T₁ spanning [-0.1354, 1.1617], T₂ spanning [-0.0992, 1.1018], T₃ spanning [-0.1956, 0.9536], and T₄ spanning [-0.2884, 0.6300]. For Schmidt 499, the predictions exhibited the expected slower diffusion characteristics with T₁ spanning [-0.1511, 0.7630], T₂ spanning [-0.1701, 0.7508], T₃ spanning [-0.1573, 0.7165], and T₄ spanning [-0.2842, 0.6555]. All predictions maintained smooth temporal evolution and physically plausible value ranges.

3.2.4 Physical Consistency of Predictions

The 3D CNN model successfully captured the essential physics distinguishing high and low Schmidt number flows. At low Sc (0.00009), the predictions correctly showed enhanced diffusion with more uniform scalar distributions, consistent with molecular diffusion dominance .

At high Sc (499), the model appropriately predicted more persistent concentration gradients and filamentary structures characteristic of advection-dominated transport. This demonstrates that the physics-informed loss function and U-Net architecture effectively learned the underlying transport mechanisms.

3.3 Computational Performance

Computational performance was highly efficient with total training time of 12.5 hours for 64 epochs, averaging 90-95 seconds per epoch. The compact 5.36 MB model size demonstrated excellent memory efficiency for handling large 3D problems. Training stability was exceptional with no observed divergence or significant oscillations throughout the entire process. The automatic learning rate reduction at epoch 64 from 5.0×10^{-5} to 2.5×10^{-5} provided appropriate optimization adjustment as convergence approached completion.

Key strengths of this training include the excellent convergence with 35-fold loss improvement, strong generalization to unseen Schmidt numbers, exceptional training stability without divergence, efficient model architecture for large 3D problems, and physically plausible predictions that follow expected diffusion patterns. The moderate validation correlation coefficients (0.69-0.77) indicate some room for improvement, while the early stopping at epoch 64 suggests potential for extended training with reduced learning rates.

The training generated comprehensive deliverables including complete model checkpoints saved every epoch, predicted fields for Schmidt numbers 9e-05 and 499, OpenFOAM-formatted fields ready for CFD simulation, and detailed training logs with epoch-by-epoch metrics.

These outputs provide immediate utility for CFD applications, parameter studies, educational demonstrations, and benchmarking for future model development. 660
661

In conclusion, the spatiotemporal 3D CNN training successfully achieved all primary 662 objectives, demonstrating robust learning of diffusion physics across varying Schmidt 663 numbers. The model shows strong predictive capabilities and generalization 664 performance, making it suitable for production use in scientific computing 665 applications.

4. Conclusion 666

This work demonstrates a successful integration of high-fidelity CFD physics with deep 667 learning to model turbulent micromixing in a manner both efficient and physically 668 interpretable. A convolutional neural network surrogate was trained on OpenFOAM-generated 669 DNS data of scalar dispersion in turbulence, focusing on capturing the Batchelor-scale 670 mixing dynamics. The resulting model reproduces the evolution of scalar concentration 671 fields with good accuracy across a wide range of Schmidt numbers, indicating that it has 672 learned the essential features of turbulent diffusion (e.g. faster spreading at low Sc, 673 filament retention at high Sc). The network's physics-informed design contributed 674 significantly to its performance: by penalizing negative concentrations and enforcing 675 smoothness, the CNN's predictions remained physically consistent and free of 676 nonphysical artifacts. Moreover, conditioning the model on the Schmidt number 677 enabled it to generalize its predictions and handle different diffusion 678 regimes within one unified framework. 679

In terms of computational benefit, the CNN surrogate provides orders-of-magnitude 679 speedup over direct CFD simulation. Once trained, the network can predict a 3D scalar field 680 in milliseconds to seconds, compared to the many hours of runtime required for a DNS on a 681 1024^3 grid. Notably, a recent study by Ouyang et al. achieved a similar acceleration ($\sim 10 \times$ 682 faster) for turbulent reactive flow simulations using a deep neural network without sacrificing 683 accuracy biblio.ugent.be . Our results align with these findings – the neural network 684 approach can bridge the traditional trade-off between accuracy and efficiency in turbulence 685 modeling biblio.ugent.be . By deploying the trained model within a CFD workflow (the 686 output is formatted to be directly usable in OpenFOAM, as in our case), one can hybridize 687 simulation and prediction, using the neural network to rapidly "fill in" fine-scale scalar 688 details in scenarios where a full DNS would be infeasible. 689

The implications for sustainable pharmaceutical process engineering are significant. Rapid 690 surrogate models for micromixing allow engineers to incorporate mixing considerations 691 early in the design of reactors and processes. For instance, the CNN model can be used to screen 692 operating conditions (stirrer speeds, feed injection patterns, etc.) to find those that ensure 693 swift micromixing, thereby suppressing side reactions and improving yield aidic.it . 694 This capability supports a move towards Digital Twins of chemical reactors – computational 695 models that can run in parallel with physical processes to provide real-time insights and 696 optimization. By reducing the computational barrier, the presented approach makes detailed 697

turbulence-level analysis accessible for routine process optimization, which was previously 698 limited to academic studies due to resource constraints. 699

Despite these advancements, there remain avenues for further refinement. One challenge 700 noted is that fully achieving Batchelor-scale resolution still requires enormous data and 701 memory; our network was trained on Kolmogorov-scale data (216^3 grid) as a stepping stone. 702 Future work will extend the training dataset to even higher resolutions as computational 703 resources grow, and explore multi-resolution techniques or neural network super-resolution 704 methods to infer fine scales from coarse inputs. Additionally, incorporating reactive chem- 705 istry into the model is a logical next step. While this study treated a passive scalar, real 706 pharmaceutical processes involve chemical reactions; embedding reaction kinetics into the 707 neural network (or coupling the CNN with a reaction model) would allow direct prediction 708 of species concentrations and yields. Physics-informed learning will be even more critical 709 in that context to ensure the network respects both mixing and reaction kinetics (e.g. by 710 enforcing that the reaction progress is diffusion-limited in high-Sc scenarios). 711

In conclusion, the presented CNN surrogate for turbulent micromixing represents a con- 712 vergence of turbulence theory, numerical simulation, and machine learning. It confirms that 713 a well-designed neural network can encapsulate complex scalar transport physics – including 714 turbulent advection and molecular diffusion at the microscale – and deliver fast, reliable 715 predictions. By applying this tool to pharmaceutical mixing problems, manufacturers can 716 achieve better control over microscale processes, leading to higher efficiency, greener chem- 717 istry (through waste minimization), and accelerated innovation in process development. The 718 approach underscores a broader paradigm in modern engineering: leveraging data-driven 719 models, guided by physical laws, to overcome traditional computational limitations and 720 drive sustainable process improvements. 721

5. Future Work

5.1 CFD Extensions

In the very near future, the plan is to use higher resolution mesh (1024^3 , 512^3) to get higher 724 accuracy of the diffusion physics. This will be done through accessing a numerical cluster 725 that enables working with such meshes. Also, the transient solver will be employed instead 726 of the steady-state solver for the same reasons (computational power limitations). These 727 improvements will allow for better resolution of the Kolmogorov and potentially Batchelor 728 scales, providing more accurate training data for neural network models. 729

5.2 Neural Network Improvements

In the future, the following improvements can be made in the neural network part: develop- 731 ment of a high-resolution database, followed by training the network with a larger number 732 of samples and higher spatial resolution. The target resolution should be $1024 \times 1024 \times 1024$ 733

or higher, as a resolution of 216^3 has been used in this project.

734

After that, additional inputs such as total time and time step can be embedded to ensure 735
that the neural network captures the correct physical behavior based on the Schmidt number 736
(Sc). If Sc is low, the refinement of the output 3D fields should be higher and vice versa. 737
Moreover, the 3D scalar fields should be stored for each time step until the total simulation 738
time specified by the user is reached. 739

Recommended next steps include extended training with reduced learning rates, valida- 740
tion on additional Schmidt numbers, architectural enhancements with residual connections, 741
multi-scale training approaches, and integration of physical constraints through diffusion 742
equation regularization. Furthermore, incorporating reactive chemistry into the model would 743
extend its applicability to real pharmaceutical manufacturing scenarios where chemical re- 744
actions occur simultaneously with mixing. 745

References

746

- [1] Jerzy Baldyga, Marta Jasińska and Michał Kotowicz. ‘Mass Transfer, Micromixing 747
and Chemical Reactions carried out in Rotor–Stator Mixers’. In: *Chemical Engineering 748
Transactions*. Vol. 57. AIDIC, 2017, pp. 1321–1326. DOI: [10.3303/CET1757221](https://doi.org/10.3303/CET1757221). 749
- [2] Jacek Pozorski and Jean-Pierre Minier. ‘On the Mixing of Passive Scalars in Turbulent 750
Flows: A Review of Physics, Modelling and DNS/LES Results’. In: *Processes* 8.11 751
(2020), p. 1379. DOI: [10.3390/pr8111379](https://doi.org/10.3390/pr8111379). 752
- [3] G. K. Batchelor. ‘Small-scale variation of convected quantities like temperature in 753
turbulent fluid. Part 1’. In: *Journal of Fluid Mechanics* 5 (1959), pp. 113–133. DOI: 754
[10.1017/S002211205900009X](https://doi.org/10.1017/S002211205900009X). 755
- [4] Guangze Li, Huangwei Zhang and Longfei Chen. ‘Large Eddy Simulation and Finite- 756
Volume Conditional Moment Closure Modelling of a Turbulent Lifted H₂/N₂ Flame’. 757
In: *International Journal of Hydrogen Energy* 46.80 (2021), pp. 40120–40142. DOI: 758
[10.1016/j.ijhydene.2021.09.209](https://doi.org/10.1016/j.ijhydene.2021.09.209). 759
- [5] Tushar Pant et al. ‘Transported PDF Modeling of Compressible Turbulent Reactive 760
Flows Using the Eulerian Monte Carlo Fields Method’. In: *Journal of Computational 761
Physics* 431 (2021), p. 110140. DOI: [10.1016/j.jcp.2020.110140](https://doi.org/10.1016/j.jcp.2020.110140). 762
- [6] George E. Karniadakis et al. ‘Physics-informed machine learning’. In: *Nature Reviews 763
Physics* 3.6 (2021), pp. 422–440. DOI: [10.1038/s42254-021-00314-5](https://doi.org/10.1038/s42254-021-00314-5). 764
- [7] Matthias Eichinger, Alexander Heinlein and Axel Klawonn. ‘Surrogate convolutional 765
neural network models for steady computational fluid dynamics simulations’. In: *Elec- 766
tronic Transactions on Numerical Analysis* 56 (2022), pp. 235–255. DOI: [10.1553/etna_vol56s235](https://doi.org/10.1553/etna_vol56s235). URL: <https://epub.oeaw.ac.at/?arp=0x003d4c21>. 768

- [8] Pratip Rana et al. ‘A scalable convolutional neural network approach to fluid flow prediction in complex environments’. In: *Scientific Reports* 14.1 (2024), p. 23080. DOI: [10.1038/s41598-024-73529-y](https://doi.org/10.1038/s41598-024-73529-y).
- [9] Yubiao Sun, Ushnish Sengupta and Matthew Juniper. ‘Physics-informed deep learning for simultaneous surrogate modeling and PDE-constrained optimization of an airfoil geometry’. In: *Computer Methods in Applied Mechanics and Engineering* 411 (2023), p. 116042. DOI: [10.1016/j.cma.2023.116042](https://doi.org/10.1016/j.cma.2023.116042).
- [10] Maziar Raissi, Hasssan Babaee and Peyman Givi. ‘Deep learning of turbulent scalar mixing’. In: *Physical Review Fluids* 4.12 (2019), p. 124501. DOI: [10.1103/PhysRevFluids.4.124501](https://doi.org/10.1103/PhysRevFluids.4.124501).
- [11] Eric Perlman et al. *Johns Hopkins Turbulence Database (JHTDB)*. <https://turbulence.idies.jhu.edu/home>. 2007.

A. Summary of Solver Modifications (scalarTransportFoam)

A.1 Purpose of the Modification

The standard scalarTransportFoam solver uses a **single, frozen velocity field** throughout the entire simulation.

For the Bayer project, the solver was modified to **support multiple, time-varying velocity fields**, enabling transient advection of a passive scalar with externally provided velocity data at each time step.

The new solver was named transientScalarTransportFoam to reflect this functionality.

A.2 Key Code-Level Changes

1. Introduced the velocity file counter and reading logic

```
// Initialize U file counter
label uCounter = 1;

// Path to next velocity field
fileName uFilePath = "all_U_files/U_" + name(uCounter);

volVectorField U_new
(
    IOobject
    (
        uFilePath,
```

```

        runTime,
803
        IOobject::MUST_READ,
804
        IOobject::AUTO_WRITE
805
),
806
mesh
807
);
808
809
if (!U_new.headerOk())
810
{
811
    FatalErrorInFunction
812
        << "Cannot read U field from file: " << uFilePath
813
        << exit(FatalError);
814
}
815
816
U = U_new;
817
uCounter++;
818

2. Recomputed flux each time step
819

phi = fvc::flux(U);
820

3. Slight modification to the control loop
821

while (simple.loop(runTime))
822
{
823
    Info<< "Time = " << runTime.timeName() << nl << endl;
824
825
    // Read next U file and update flux
826
    ...
827
}
828

```

B. MATLAB Script Modifications for HDF5 to Open-FOAM Conversion

B.1 Purpose of the Modification

The original MATLAB script was designed to **visualize scalar fields** stored in .h5 format. 832
 For the Bayer project, it was extended to **convert these .h5 files into OpenFOAM- 833
 compatible formats**, enabling seamless use of external velocity field data in the transient 834
 scalar transport solver. 835

B.2 Key Code-Level Changes

836

1. Added batch processing of multiple .h5 files

837

```
fileList = dir(fullfile(dataPath, '*.h5'));
nFiles = length(fileList);
for i = 1:nFiles
    fileName = fullfile(dataPath, fileList(i).name);
    data = h5read(fileName, '/velocity');
    ...
end
```

838

839

840

841

```
    data = h5read(fileName, '/velocity');
```

842

843

844

```
    ...
```

845

2. Conversion of velocity components to OpenFOAM format

846

```
Ux = squeeze(data(1,:,:,:));
Uy = squeeze(data(2,:,:,:));
Uz = squeeze(data(3,:,:,:));
outputFile = fullfile(outputDir, sprintf('U_%d', i));
writeOpenFOAMVector(outputFile, Ux, Uy, Uz);
```

847

848

849

850

851

852

3. Introduced a custom writing function

853

```
function writeOpenFOAMVector(filename, Ux, Uy, Uz)
fid = fopen(filename, 'w');
fprintf(fid, "FoamFile\n{\n version 2.0;\n format ascii;\n class volVectorField;\n}\n");
...
for cell = 1:numel(Ux)
    fprintf(fid, "(%f %f %f)\n", Ux(cell), Uy(cell), Uz(cell));
end
fclose(fid);
end
```

854

855

856

857

858

859

860

861

862

4. Automated mesh/coordinate mapping

863

```
[x, y, z] = ndgrid(xCoords, yCoords, zCoords);
% optional visualization
quiver3(x(:, ), y(:, ), z(:, ), Ux(:, ), Uy(:, ), Uz(:, ));
```

864

865

866

C. 3D CNN Implementation Code

867

The 3D CNN code `tf_test.py` implements the spatiotemporal physics-informed neural network for predicting turbulent scalar mixing. The complete implementation includes:

868

869

- Custom PhysicsAwareLoss class incorporating non-negativity and total variation penalties 870
- 871
- 3D U-Net architecture with Schmidt number conditioning 872
- 873
- Data preprocessing pipeline for OpenFOAM field files 873
- 874
- Training loop with early stopping and learning rate scheduling 874
- 875
- Output generation in OpenFOAM-compatible format 875

```

# Physics-informed loss function
class PhysicsAwareLoss(tf.keras.losses.Loss):
    def __init__(self, lambda_tv=0.01, lambda_neg=1.0):
        super().__init__()
        self.lambda_tv = lambda_tv
        self.lambda_neg = lambda_neg

    def call(self, y_true, y_pred):
        # Mean squared error
        mse = tf.reduce_mean(tf.square(y_true - y_pred))

        # Non-negativity penalty
        neg_penalty = tf.reduce_mean(tf.nn.relu(-y_pred))

        # Total variation (smoothness) in 3D
        tv_x = tf.reduce_mean(tf.abs(y_pred[:, 1:, :, :, :] - y_pred[:, :-1, :, :, :]))
        tv_y = tf.reduce_mean(tf.abs(y_pred[:, :, 1:, :, :] - y_pred[:, :, :-1, :, :]))
        tv_z = tf.reduce_mean(tf.abs(y_pred[:, :, :, 1:, :] - y_pred[:, :, :, :-1, :]))
        tv_penalty = tv_x + tv_y + tv_z

        # Combined loss
        total_loss = mse + self.lambda_neg * neg_penalty + self.lambda_tv * tv_penalty
        return total_loss

# 3D U-Net model with Schmidt number conditioning
# (Full model definition as shown in Listing 3)
# ...

# Training configuration
model.compile(

```

```

optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4),
loss=PhysicsAwareLoss(lambda_tv=0.01, lambda_neg=1.0),
metrics=['mae', 'mse']
)

# Callbacks for training stability
callbacks = [
    tf.keras.callbacks.EarlyStopping(patience=10, restore_best_weights=True),
    tf.keras.callbacks.ReduceLROnPlateau(factor=0.5, patience=5),
    tf.keras.callbacks.ModelCheckpoint('best_model.h5', save_best_only=True)
]

# Train the model
history = model.fit(
    [initial_fields, schmidt_numbers],
    target_fields,
    batch_size=1,
    epochs=100,
    validation_split=0.2,
    callbacks=callbacks
)

```

Listing 4: Spatiotemporal 3D CNN Implementation (Excerpt)

The full code is available in the project repository and includes additional functionality 933
for data augmentation, field normalization, and OpenFOAM output formatting. 934